

## Section 2: Yelp Data Challenge - NLP

Yiting Luo | Data Science Applied Research - 2

May 2018

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv('dataset/last_2_years_restaurant_reviews.csv')
```

```
In [3]: df.head()
```

Out[3]:

	business_id	name	categories	avg_stars	cool	date	funny	
0	-9e10NYQuAa-CB_Rrw7Tw	Delmonico Steakhouse	['Cajun/Creole', 'Steakhouses', 'Restaurants']	4.0	0	2016-03-31	0	6SgvNWJI
1	-9e10NYQuAa-CB_Rrw7Tw	Delmonico Steakhouse	['Cajun/Creole', 'Steakhouses', 'Restaurants']	4.0	0	2015-06-29	0	iwx6s6yQ
2	-9e10NYQuAa-CB_Rrw7Tw	Delmonico Steakhouse	['Cajun/Creole', 'Steakhouses', 'Restaurants']	4.0	0	2015-03-16	0	UVUMu_b
3	-9e10NYQuAa-CB_Rrw7Tw	Delmonico Steakhouse	['Cajun/Creole', 'Steakhouses', 'Restaurants']	4.0	0	2016-02-10	0	UxFpgng8
4	-9e10NYQuAa-CB_Rrw7Tw	Delmonico Steakhouse	['Cajun/Creole', 'Steakhouses', 'Restaurants']	4.0	0	2017-02-14	0	Xp3ppynE

## Define your feature variables, here is the text of the review

```
In [4]: # Take the values of the column that contains review text data, save it to a variable named "documents"
documents = df['text'].values # Make results to Numpy array
```

```
In [5]: # inspect your documents, e.g. check the size, take a peek at elements of the numpy array
documents.dtype, documents.shape
```

Out[5]: (dtype('O'), (515752,))

```
In [6]: documents[3]
```

```
Out[6]: 'Truly Fantastic! Best Steak ever. Service was awesome and timely.
They knew right when you needed something. The root beer float for d
essert was great.'
```

## Define target variable (any categorical variable that may be meaningful)

I am interested in perfect (5 stars) and imperfect (1-4 stars) rating

```
In [7]: # Make a column and take the values, save to a variable named "target"
df['favorable'] = (df['stars'] > 4)
```

```
In [8]: target = df['favorable'].values
```

```
In [9]: target[:10]
```

```
Out[9]: array([ True, False,  True,  True,  True, False,  True,  True,  True
,
          False])
```

Look at the statistic of the target variable

```
In [10]: # To be implemented
target.mean(), target.std(), documents.shape, target.shape
```

```
Out[10]: (0.46397299477268145, 0.49870036584541505, (515752,), (515752,))
```

## Let's create training dataset and test dataset

```
In [11]: from sklearn.cross_validation import train_test_split
```

```
/Users/luoyiting/anaconda/lib/python3.5/site-packages/sklearn/cross_
validation.py:44: DeprecationWarning: This module was deprecated in
version 0.18 in favor of the model_selection module into which all t
he refactored classes and functions are moved. Also note that the in
terface of the new CV iterators are different from that of this modu
le. This module will be removed in 0.20.
  "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [12]: # Documents is X, target is y
# Now split the data to training set and test set
```

```
In [13]: # Split to documents_train, documents_test, target_train, target_test
documents_train, documents_test, target_train, target_test = train_test_split(
    documents, target, test_size = 0.8, random_state = 42
)
```

## Let's get NLP representation of the documents

```
In [14]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [15]: # Create TfidfVectorizer, and name it vectorizer
vectorizer = TfidfVectorizer(stop_words = 'english', max_features = 5000)
```

```
In [16]: # Train the model with training data
vectors_train = vectorizer.fit_transform(documents_train).toarray() #
toarray avoid sparse matrix
```

```
/Users/luoyiting/anaconda/lib/python3.5/site-packages/sklearn/feature_extraction/text.py:1059: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
    if hasattr(X, 'dtype') and np.issubdtype(X.dtype, np.float):
```

```
In [17]: # Get the vocab of tfidf
words = vectorizer.get_feature_names()
```

```
In [18]: vectors_train.shape
```

```
Out[18]: (103150, 5000)
```

```
In [19]: # Use the trained model to transform your test data
vectors_test = vectorizer.transform(documents_test).toarray()
```

```
/Users/luoyiting/anaconda/lib/python3.5/site-packages/sklearn/feature_extraction/text.py:1059: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
    if hasattr(X, 'dtype') and np.issubdtype(X.dtype, np.float):
```

## Similar review search engine

```
In [20]: import numpy as np
```

```
def get_top_values(lst, n, labels):  
    '''  
    INPUT: LIST, INTEGER, LIST  
    OUTPUT: LIST  
  
    Given a list of values, find the indices with the highest n values  
    .  
    Return the labels for each of these indices.  
  
    e.g.  
    lst = [7, 3, 2, 4, 1]  
    n = 2  
    labels = ["cat", "dog", "mouse", "pig", "rabbit"]  
    output: ["cat", "pig"]  
    '''  
    return [labels[i] for i in np.argsort(lst)[::-1][:n]] # np.argsort  
t by default sorts values in ascending order  
  
def get_bottom_values(lst, n, labels):  
    '''  
    INPUT: LIST, INTEGER, LIST  
    OUTPUT: LIST  
  
    Given a list of values, find the indices with the lowest n values.  
    Return the labels for each of these indices.  
  
    e.g.  
    lst = [7, 3, 2, 4, 1]  
    n = 2  
    labels = ["cat", "dog", "mouse", "pig", "rabbit"]  
    output: ["mouse", "rabbit"]  
    '''  
    return [labels[i] for i in np.argsort(lst)[:n]]
```

```
In [21]: # Let's use cosine similarity  
from sklearn.metrics.pairwise import cosine_similarity
```

```
In [22]: # Draw an arbitrary review from test (unseen in training) documents
some_random_number = 42
search_query = documents_test[some_random_number]
search_queries = [search_query]
print (search_query)
print (search_queries)
```

Great food and great price, the only thing is the waiting time. During lunch and dinner hours, they are extremely busy. Call in ahead to get your order and pick up for faster service :)

['Great food and great price, the only thing is the waiting time. During lunch and dinner hours, they are extremely busy. Call in ahead to get your order and pick up for faster service :)']

```
In [23]: # Transform the drawn review(s) to vector(s)
vector_search_queries = vectorizer.transform(search_queries).toarray()

/Users/luoyiting/anaconda/lib/python3.5/site-packages/sklearn/feature_extraction/text.py:1059: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
    if hasattr(X, 'dtype') and np.issubdtype(X.dtype, np.float):
```

```
In [24]: # Calculate the similarity score(s) between vector(s) and training vectors
similarity_score = cosine_similarity(vector_search_queries, vectors_train)
```

```
In [25]: similarity_score.shape
```

```
Out[25]: (1, 103150)
```

```
In [26]: # Let's find top 5 similar reviews
n = 5
returned_reviews = get_top_values(similarity_score[0], n, documents_train)
```

```
In [27]: print ('Our search query:')
print (search_queries[0])
```

Our search query:

Great food and great price, the only thing is the waiting time. During lunch and dinner hours, they are extremely busy. Call in ahead to get your order and pick up for faster service :)

```
In [28]: print ('Most %s similar reviews:' % n)
        for review in returned_reviews:
            print (review )
```

Most 5 similar reviews:

Great buffet and a great price. Ive eaten here 4 times. It was great each time. It can get busy though so i would suggest call ahead to find out the wait time

Well I'm updating my last review. As previously stated gave them 2 stars because the food is somewhat good. I don't like dealing with the depressed people inside so I decided moving forward I would place the order online and pick up.

When I arrived a few minutes after my pick up time they told me sorry your order is not ready, it will be faster if you stand in line. How is it faster if I stand in line when have had my order for 45 minutes?! They told me it would be an additional 20 minutes to make my order but they could magically make it faster if I stood in line in 10 min before ordering the exact same thing they had on the receipt in front of them. No sense in coming back. I have tried time and time again, and every time it gets worse. This will be my last trip to place I at one point in my life came to weekly. So with this, so long cafe rio.

For lunch or dinner, get the BBQ burger. One of the best I ever had. Great food, service, price and location.

I had a great time today going to the barbecue restaurant. It is the first time of that I have been there, and the service is excellent much faster than I expected, and the food was delicious!

Great place, Great Food, EXCELLENT CUSTOMER SERVICE!! We came in with 8 ppl they fixed us up a table and served us faster than McDonalds serves Fries! Great quality food... Everything is Great Great Great!!! If you haven't ate here its a MUST!!

The search engine makes sense.

## Classifying positive/negative review

### Naive-Bayes Classifier

```
In [29]: # Build a Naive-Bayes Classifier

        from sklearn.naive_bayes import MultinomialNB

        model_nb = MultinomialNB()
        model_nb.fit(vectors_train, target_train)
```

```
Out[29]: MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)
```

```
In [30]: # Get score for training set  
model_nb.score(vectors_train, target_train) #accuracy
```

```
Out[30]: 0.8092195831313621
```

```
In [31]: # Get score for test set  
model_nb.score(vectors_test, target_test)
```

```
Out[31]: 0.8031250454433085
```

## Logistic Regression Classifier

```
In [32]: # Build a Logistic Regression Classifier  
  
from sklearn.linear_model import LogisticRegression  
  
model_lrc = LogisticRegression()  
model_lrc.fit(vectors_train, target_train)
```

```
Out[32]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_interce  
pt=True,  
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jo  
bs=1,  
                             penalty='l2', random_state=None, solver='liblinear', tol=0  
.0001,  
                             verbose=0, warm_start=False)
```

```
In [33]: # Get score for training set  
model_lrc.score(vectors_train, target_train) #accuracy
```

```
Out[33]: 0.8433446437227339
```

```
In [34]: # Get score for test set  
model_lrc.score(vectors_test, target_test)
```

```
Out[34]: 0.8256940102083848
```

**Key features(words) that make the positive prediction?**



```
In [35]: # find it out by ranking
n = 20
get_top_values(model_lrc.coef_[0], n, words)
```

```
Out[35]: ['amazing',
          'best',
          'incredible',
          'delicious',
          'awesome',
          'thank',
          'perfection',
          'perfect',
          'phenomenal',
          'fantastic',
          'favorite',
          'great',
          'perfectly',
          'excellent',
          'highly',
          'die',
          'heaven',
          'gem',
          'love',
          'bomb']
```

**Q: Key features(words) that make the negative prediction?**

```
In [36]: # find it out by ranking
n = 20
get_bottom_values(model_lrc.coef_[0], n, words)
```

```
Out[36]: ['worst',
          'ok',
          'horrible',
          'rude',
          'bland',
          'terrible',
          'okay',
          'mediocre',
          'slow',
          'disappointing',
          'dry',
          'meh',
          'lacking',
          'unfortunately',
          'overpriced',
          'wasn',
          'awful',
          'average',
          'poor',
          'decent']
```

## Random Forest Classifier

```
In [37]: # Build a Random Forest Classifier

from sklearn.ensemble import RandomForestClassifier

model_rfc = RandomForestClassifier(max_depth = None,
                                   n_estimators = 5,
                                   min_samples_leaf = 10)
model_rfc.fit(vectors_train, target_train)
```

```
Out[37]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion=
'gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None
,
                                min_impurity_split=1e-07, min_samples_leaf=10,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                n_estimators=5, n_jobs=1, oob_score=False, random_state=
None,
                                verbose=0, warm_start=False)
```

```
In [38]: # Get score for training set
model_rfc.score(vectors_train, target_train)
```

```
Out[38]: 0.8138826951042172
```

```
In [39]: # Get score for test set  
model_rfc.score(vectors_test, target_test)
```

```
Out[39]: 0.7711305325713399
```

### What features (words) are important by inspecting the RFC model?

```
In [41]: n = 20  
get_top_values(model_rfc.feature_importances_, n, words)
```

```
Out[41]: ['amazing',  
          'great',  
          'love',  
          'best',  
          'awesome',  
          'delicious',  
          'bad',  
          'like',  
          'friendly',  
          'didn',  
          'horrible',  
          'favorite',  
          'highly',  
          'excellent',  
          'perfect',  
          'definitely',  
          'wasn',  
          'average',  
          'thank',  
          'worst']
```

## Use cross validation to evaluate classifiers

sklearn cross validation ([http://scikit-learn.org/stable/modules/cross\\_validation.html](http://scikit-learn.org/stable/modules/cross_validation.html))

```
In [42]: from sklearn.model_selection import cross_val_score  
  
cv_scores = cross_val_score(model_lrc,  
                             vectors_train,  
                             target_train,  
                             cv = 5,  
                             scoring = "accuracy"  
                             )
```

```
In [44]: cv_scores.mean(), cv_scores.std()
```

```
Out[44]: (0.826136643215901, 0.003967095258931436)
```

# Use grid search to find best predictable classifier

[sklearn grid search tutorial \(with cross validation\) \(http://scikit-learn.org/stable/modules/grid\\_search.html#grid-search\)](http://scikit-learn.org/stable/modules/grid_search.html#grid-search)

[sklearn grid search documentation \(with cross validation\) \(http://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.GridSearchCV.html#sklearn.model\\_selection.GridSearchCV\)](http://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html#sklearn.model_selection.GridSearchCV)

```
In [45]: # Tune Logistic Regression Regularization parameter C and different penalty
         from sklearn.model_selection import GridSearchCV
         from sklearn.metrics import classification_report

         param_grid = [{'penalty':['l1'], 'C':[0.1, 100]},
                        {'penalty':['l2'], 'C':[0.1, 100]}]

         scores = ['accuracy']

         for score in scores:
             print("# Tuning hyper-parameters for %s" % score + "\n\n")
             clf = GridSearchCV(LogisticRegression(),
                                param_grid,
                                cv=5,
                                scoring=score)
             clf.fit(vectors_train[:500,:], target_train[:500])
             print("Best parameters set found on development set:\n\n")
             print(clf.best_params_)
             print("\nGrid scores on development set:\n\n")
             means = clf.cv_results_['mean_test_score']
             stds = clf.cv_results_['std_test_score']
             for mean, std, params in zip(means, stds, clf.cv_results_['params']):
                 print("%0.3f (+/-%0.03f) for %r"
                       % (mean, std * 2, params))

             print("\nDetailed classification report:\n")
             print("The model is trained on the full development set.")
             print("The scores are computed on the full evaluation set.")
             print("\n")
             y_true, y_pred = target_test, clf.predict(vectors_test)
             print(classification_report(y_true, y_pred))
             print("\n")
```

```
# Tuning hyper-parameters for accuracy
```

```
Best parameters set found on development set:
```

```
{'penalty': 'l2', 'C': 100}
```

```
Grid scores on development set:
```

```
0.554 (+/-0.004) for {'penalty': 'l1', 'C': 0.1}  
0.718 (+/-0.083) for {'penalty': 'l1', 'C': 100}  
0.552 (+/-0.012) for {'penalty': 'l2', 'C': 0.1}  
0.752 (+/-0.073) for {'penalty': 'l2', 'C': 100}
```

```
Detailed classification report:
```

```
The model is trained on the full development set.  
The scores are computed on the full evaluation set.
```

	precision	recall	f1-score	support
False	0.75	0.78	0.76	221181
True	0.73	0.69	0.71	191421
avg / total	0.74	0.74	0.74	412602

```
In [ ]:
```